

# Predict monthly Car Sales

February 14, 2022

```
[217]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.model_selection import train_test_split, RandomizedSearchCV,
    GridSearchCV
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.multioutput import MultiOutputRegressor
from xgboost import XGBRegressor
from sklearn.preprocessing import MinMaxScaler, StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.neural_network import MLPRegressor
import tensorflow as tf
import keras
from warnings import simplefilter
simplefilter("ignore")
# Set Matplotlib defaults
plt.style.use("seaborn-whitegrid")
plt.rc("figure", autolayout=True, figsize=(11, 5))
plt.rc("axes", labelweight="bold", labelsiz="large", titleweight="bold",
    titlesize=16, titlepad=10)
plot_params = dict(color="0.75", style=".-", markeredgecolor="0.25",
    markerfacecolor="0.25", legend=False)
%config InlineBackend.figure_format = 'retina'
np.random.seed(42)
```

```
[45]: # define train test split function
def train_test_datasets(df, x_len=12, y_len=1, test_loops=12):
    D = df.values
    rows, periods = D.shape

    # Training set creation
    loops = periods + 1 - x_len - y_len
    train = []
```

```

for col in range(loops):
    train.append(D[:, col:col+x_len+y_len])
train = np.vstack(train)
X_train, y_train = np.split(train, [-y_len], axis=1)

# Test set creation
if test_loops > 0:
    X_train, X_test = np.split(X_train, [-rows*test_loops], axis=0)
    y_train, y_test = np.split(y_train, [-rows*test_loops], axis=0)
else: # No test set: X_test is used to generate the future forecast
    X_test = D[:, -x_len:]
    y_test = np.full((X_test.shape[0], y_len), np.nan) #Dummy value

# Formatting required for scikit-learn
if y_len == 1:
    y_train = y_train.ravel()
    y_test = y_test.ravel()

return X_train, y_train, X_test, y_test

# define score metric function
def kpi(y_train, y_train_pred, y_test, y_test_pred, name=''):
    df = pd.DataFrame(columns = ['MAE', 'RMSE', 'Bias', 'MAE_pct', 'RMSE_pct', 'r2_score'], index=['Train', 'Test'])
    df.index.name = name
    df.loc['Train', 'MAE_pct'] = 100*np.mean(abs(y_train - y_train_pred))/np.
    mean(y_train)
    df.loc['Train', 'RMSE_pct'] = 100*np.sqrt(np.mean((y_train -
    y_train_pred)**2))/np.mean(y_train)
    df.loc['Train', 'Bias'] = 100*np.mean((y_train - y_train_pred))/np.
    mean(y_train)
    df.loc['Train', 'r2_score'] = r2_score(y_train, y_train_pred)
    df.loc['Train', 'MAE'] = mean_absolute_error(y_train, y_train_pred)
    df.loc['Train', 'RMSE'] = mean_squared_error(y_train, y_train_pred,
    squared=False)
    df.loc['Test', 'MAE_pct'] = 100*np.mean(abs(y_test - y_test_pred))/np.
    mean(y_test)
    df.loc['Test', 'RMSE_pct'] = 100*np.sqrt(np.mean((y_test - y_test_pred)**2))/
    np.mean(y_test)
    df.loc['Test', 'Bias'] = 100*np.mean((y_test - y_test_pred))/np.mean(y_test)
    df.loc['Test', 'r2_score'] = r2_score(y_test, y_test_pred)
    df.loc['Test', 'MAE'] = mean_absolute_error(y_test, y_test_pred)
    df.loc['Test', 'RMSE'] = mean_squared_error(y_test, y_test_pred,
    squared=False)
    df = df.astype(float).round(2) #Round number for display
    print(df)

```

```
[442]: data = pd.read_csv('norway_new_car_sales_by_make.csv')
data.tail()
```

```
[442]:      Year  Month      Make  Quantity  Pct
3947  2016      1      Smart         6  0.1
3948  2016      1       Fiat         5  0.0
3949  2016      1       Jeep         4  0.0
3950  2016      1      Dacia         1  0.0
3951  2016      1  Maserati         1  0.0
```

```
[443]: data['Period'] = data.Year.astype(str) + '-' + data.Month.astype(str).str.
      ↪zfill(2)
```

```
[444]: data.head()
```

```
[444]:      Year  Month      Make  Quantity  Pct  Period
0  2007      1      Toyota      2884  22.7  2007-01
1  2007      1  Volkswagen      2521  19.9  2007-01
2  2007      1     Peugeot      1029   8.1  2007-01
3  2007      1        Ford       870   6.9  2007-01
4  2007      1        Volvo       693   5.5  2007-01
```

```
[445]: data.isnull().sum()
```

```
[445]: Year          0
Month          0
Make          10
Quantity       0
Pct           0
Period        0
dtype: int64
```

```
[446]: data[data.Make.isnull()] # There is months with out cars sales we will fill 0
```

```
[446]:      Year  Month Make  Quantity  Pct  Period
37   2007      1  NaN         1  0.0  2007-01
112  2007      3  NaN         1  0.0  2007-03
265  2007      7  NaN         1  0.0  2007-07
419  2007     11  NaN         1  0.0  2007-11
1256 2009      9  NaN         4  0.0  2009-09
1294 2009     10  NaN         4  0.0  2009-10
2399 2012      4  NaN         1  0.0  2012-04
2478 2012      6  NaN         1  0.0  2012-06
2517 2012      7  NaN         1  0.0  2012-07
3013 2013      9  NaN         1  0.0  2013-09
```

```
[447]: df = pd.pivot_table(data=data, values='Quantity', index='Make',
    ↪columns='Period', aggfunc='sum', fill_value=0)
df.head()
```

```
[447]: Period      2007-01  2007-02  2007-03  2007-04  2007-05  2007-06  2007-07  \
Make
Alfa Romeo      16        9       21       20       17       21       14
Aston Martin     0        0        1        0        4        3        3
Audi            599      498      682      556      630      498      562
BMW             352      335      365      360      431      477      403
Bentley          0        0        0        0        0        1        0

Period      2007-08  2007-09  2007-10  ...  2015-04  2015-05  2015-06  \
Make
Alfa Romeo      12       15       10  ...        3        3        3
Aston Martin     0        0        0  ...        2        2        0
Audi            590      393      554  ...      665      585      640
BMW             348      271      562  ...      733      693      849
Bentley          0        0        0  ...        0        0        0

Period      2015-07  2015-08  2015-09  2015-10  2015-11  2015-12  2016-01
Make
Alfa Romeo        9        4        5        0        3        3        0
Aston Martin       0        0        0        0        0        0        0
Audi             754      541      494      549      592      496      559
BMW              617      860      777     1010      934     1024     1089
Bentley            0        0        0        0        0        0        0

[5 rows x 109 columns]
```

```
[448]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 65 entries, Alfa Romeo to Westfield
Columns: 109 entries, 2007-01 to 2016-01
dtypes: int64(109)
memory usage: 55.9+ KB
```

```
[449]: df.shape
```

```
[449]: (65, 109)
```

## 0.1 Linear Regression

```
[450]: # Linear Regression as first model to compare

X_train, y_train, X_test, y_test = train_test_datasets(df, x_len=12, y_len=1,
↳test_loops=12)

reg = Pipeline([('scaler', MinMaxScaler()), ('LinearRegression',
↳LinearRegression())]) # Create a linear regression object

reg.fit(X_train, y_train) # Fit it to the training data

# Create two predictions for the training and test sets
y_train_pred = reg.predict(X_train)
y_test_pred = reg.predict(X_test)

kpi(y_train, y_train_pred, y_test, y_test_pred, name='Regression')
```

	MAE	RMSE	Bias	MAE_pct	RMSE_pct	r2_score
Regression						
Train	29.50	73.02	-0.0	17.91	44.32	0.95
Test	33.56	80.90	0.9	17.32	41.75	0.95

## 0.2 Random Forest

```
[209]: # use Random Forest with default parameters

X_train, y_train, X_test, y_test = train_test_datasets(df, x_len=12, y_len=1,
↳test_loops=12)

forest = Pipeline([('scaler', MinMaxScaler()), ('Forest',
↳RandomForestRegressor())])

forest.fit(X_train, y_train)

y_train_pred = forest.predict(X_train)
y_test_pred = forest.predict(X_test)

kpi(y_train, y_train_pred, y_test, y_test_pred, name='Forest')
```

	MAE	RMSE	Bias	MAE_pct	RMSE_pct	r2_score
Forest						
Train	11.23	28.60	-0.14	6.67	16.99	0.99
Test	35.00	89.76	1.97	17.43	44.69	0.95

```
[215]: X_train, y_train, X_test, y_test = train_test_datasets(df, x_len=12, y_len=1,
↳test_loops=12)
```

```

pipe_grid = Pipeline([('scaler', MinMaxScaler()), ('Forestb',
↳ RandomForestRegressor(n_jobs=1))])

param_dist = {'Forestb__n_estimators': list(range(100, 500, 100)),
              'Forestb__max_depth': list(range(1, 20)) + [None],
              'Forestb__min_samples_split': range(2, 20),
              'Forestb__min_samples_leaf': range(1, 20),
              'Forestb__max_features': [.1, .2, .3, .4, .5, .6, .7, .8, .9, 1]
↳+ ['auto'],
              'Forestb__bootstrap': [True],
              'Forestb__max_samples': [.1, .2, .3, .4, .5, .6, .7, .8, .9, 1]}

forest_cv = RandomizedSearchCV(pipe_grid, param_dist, cv=6, n_jobs=-1,
↳ verbose=1, n_iter=400)

forest_cv.fit(X_train, y_train)

print('Tuned Forest Parameters:', forest_cv.best_params_)

```

Fitting 6 folds for each of 400 candidates, totalling 2400 fits  
Tuned Forest Parameters: {'Forestb\_\_n\_estimators': 100,  
'Forestb\_\_min\_samples\_split': 7, 'Forestb\_\_min\_samples\_leaf': 6,  
'Forestb\_\_max\_samples': 0.5, 'Forestb\_\_max\_features': 0.4, 'Forestb\_\_max\_depth':  
19, 'Forestb\_\_bootstrap': True}

```

[216]: X_train, y_train, X_test, y_test = train_test_datasets(df, x_len=12, y_len=1,
↳ test_loops=12)

pipe_grid_f = Pipeline([('scaler', MinMaxScaler()), ('Forestf',
↳ RandomForestRegressor(n_jobs=1))])

param_dist_f = {'Forestf__n_estimators': [100],
                'Forestf__max_depth': list(range(17, 22)) + [None],
                'Forestf__min_samples_split': range(5, 10),
                'Forestf__min_samples_leaf': range(4, 9),
                'Forestf__max_features': [.35, .4, .45],
                'Forestf__bootstrap': [True],
                'Forestf__max_samples': [.45, .5, .55]}

forest_cv_f = GridSearchCV(pipe_grid_f, param_dist_f, n_jobs=-1, verbose=1)

forest_cv_f.fit(X_train, y_train)

print('Tuned Forest Parameters:', forest_cv_f.best_params_)

```

Fitting 5 folds for each of 1350 candidates, totalling 6750 fits

Tuned Forest Parameters: {'Forestf\_\_bootstrap': True, 'Forestf\_\_max\_depth': 19, 'Forestf\_\_max\_features': 0.4, 'Forestf\_\_max\_samples': 0.5, 'Forestf\_\_min\_samples\_leaf': 5, 'Forestf\_\_min\_samples\_split': 5, 'Forestf\_\_n\_estimators': 100}

```
[457]: X_train, y_train, X_test, y_test = train_test_datasets(df, x_len=12, y_len=1,
↳test_loops=20)


forest_final = Pipeline([('scaler', MinMaxScaler()), ('ForestFinal',
↳RandomForestRegressor(n_estimators=100,
↳
↳min_samples_split=5,
↳
↳min_samples_leaf=5,
↳
↳max_samples=0.5,
↳
↳max_features=0.4,
↳
↳max_depth=19,
↳
↳bootstrap=True)))]

forest_final.fit(X_train, y_train)

y_train_pred = forest_final.predict(X_train)
y_test_pred = forest_final.predict(X_test)
kpi(y_train, y_train_pred, y_test, y_test_pred, name='forest_final')
```

	MAE	RMSE	Bias	MAE_pct	RMSE_pct	r2_score
forest_final						
Train	23.53	60.29	-0.07	14.42	36.96	0.97
Test	34.27	94.47	1.89	18.19	50.14	0.94

```
[452]: cols = X_train.shape[1]
features = [f'M-{cols-col}' for col in range(cols)]
data = forest_final.steps[1][1].feature_importances_.reshape(-1,1)
imp = pd.DataFrame(data=data, index=features, columns=['Forest'])
imp.plot(kind='bar'); # the most important months is the first previous three
↳months
```



```
[456]: # prediction

X_train, y_train, X_test, y_test = train_test_datasets(df, x_len=12, y_len=1,
↳test_loops=0)

forest_final = Pipeline([('scaler', MinMaxScaler()), ('ForestFinal',
↳RandomForestRegressor(n_estimators=100,

↳ min_samples_split=5,

↳ min_samples_leaf=5,

↳ max_samples=0.5,

↳ max_features=0.4,

↳ max_depth=19,

↳ bootstrap=True))])

forest_final.fit(X_train, y_train)

forecast = pd.DataFrame(data=forest_final.predict(X_test), index=df.index,
↳columns=['Next_Sales'])
forecast.tail(10)
```

```
[456]:          Next_Sales
Make
Subaru      273.765626
```



Suzuki	246.159222
Tata	0.035983
Tazzari	0.035983
Tesla	238.629376
Think	0.035983
Toyota	1279.202421
Volkswagen	1980.871566
Volvo	754.306872
Westfield	0.035983

### 0.3 sklearn neural network

```
[239]: X_train, y_train, X_test, y_test = train_test_datasets(df, x_len=12, y_len=1,
    ↪test_loops=12)

param_fixed = {'NN__activation': 'relu',
               'NN__solver': 'adam',
               'NN__early_stopping': True,
               'NN__n_iter_no_change': 50,
               'NN__validation_fraction': .1,
               'NN__tol': .0001}

NN_b = Pipeline([('scaler', MinMaxScaler()),
                 ('NN', MLPRegressor(hidden_layer_sizes=(20, 20),
    ↪**param_fixed, verbose=True))])

NN_b.fit(X_train, y_train)

y_train_pred = NN_b.predict(X_train)
y_test_pred = NN_b.predict(X_test)

kpi(y_train, y_train_pred, y_test, y_test_pred, name='NN_b')
```

```
Iteration 1, loss = 16971.36834111
Validation score: 0.920519
Iteration 2, loss = 5110.23720794
Validation score: 0.931789
Iteration 3, loss = 4565.24989091
Validation score: 0.939719
Iteration 4, loss = 3852.89505267
Validation score: 0.947731
Iteration 5, loss = 3449.10170482
Validation score: 0.952096
Iteration 6, loss = 3265.13428419
Validation score: 0.953008
Iteration 7, loss = 3209.02490989
Validation score: 0.954035
```

Iteration 8, loss = 3085.55584211  
Validation score: 0.951573  
Iteration 9, loss = 3064.69334411  
Validation score: 0.954850  
Iteration 10, loss = 3008.82441924  
Validation score: 0.953426  
Iteration 11, loss = 2999.89364823  
Validation score: 0.954806  
Iteration 12, loss = 2935.42363800  
Validation score: 0.950930  
Iteration 13, loss = 2945.29979431  
Validation score: 0.952048  
Iteration 14, loss = 2892.57673994  
Validation score: 0.954456  
Iteration 15, loss = 2898.20545926  
Validation score: 0.953670  
Iteration 16, loss = 2840.81105990  
Validation score: 0.956535  
Iteration 17, loss = 2853.17051733  
Validation score: 0.955087  
Iteration 18, loss = 2811.50052428  
Validation score: 0.954550  
Iteration 19, loss = 2859.53208914  
Validation score: 0.957451  
Iteration 20, loss = 2801.37358361  
Validation score: 0.956494  
Iteration 21, loss = 2826.67468237  
Validation score: 0.953650  
Iteration 22, loss = 2800.97152429  
Validation score: 0.957574  
Iteration 23, loss = 2774.53435708  
Validation score: 0.957465  
Iteration 24, loss = 2786.25617047  
Validation score: 0.957286  
Iteration 25, loss = 2759.36054757  
Validation score: 0.957647  
Iteration 26, loss = 2768.49175650  
Validation score: 0.956920  
Iteration 27, loss = 2775.32711925  
Validation score: 0.957906  
Iteration 28, loss = 2737.64666304  
Validation score: 0.957604  
Iteration 29, loss = 2740.10605118  
Validation score: 0.958141  
Iteration 30, loss = 2723.86415606  
Validation score: 0.957617  
Iteration 31, loss = 2734.11520275  
Validation score: 0.956248

Iteration 32, loss = 2718.93903727  
Validation score: 0.957583  
Iteration 33, loss = 2708.15187273  
Validation score: 0.957827  
Iteration 34, loss = 2794.12130042  
Validation score: 0.956448  
Iteration 35, loss = 2717.14258181  
Validation score: 0.958125  
Iteration 36, loss = 2729.33100488  
Validation score: 0.956359  
Iteration 37, loss = 2807.90123038  
Validation score: 0.954742  
Iteration 38, loss = 2680.45248503  
Validation score: 0.954734  
Iteration 39, loss = 2678.08721670  
Validation score: 0.956289  
Iteration 40, loss = 2696.67853147  
Validation score: 0.958213  
Iteration 41, loss = 2660.16777892  
Validation score: 0.958391  
Iteration 42, loss = 2699.44483095  
Validation score: 0.957891  
Iteration 43, loss = 2671.49539269  
Validation score: 0.955042  
Iteration 44, loss = 2694.94566715  
Validation score: 0.958042  
Iteration 45, loss = 2657.01530109  
Validation score: 0.958134  
Iteration 46, loss = 2648.74621629  
Validation score: 0.957668  
Iteration 47, loss = 2637.09705193  
Validation score: 0.958059  
Iteration 48, loss = 2641.89330288  
Validation score: 0.958081  
Iteration 49, loss = 2639.54967327  
Validation score: 0.957831  
Iteration 50, loss = 2627.73160091  
Validation score: 0.953220  
Iteration 51, loss = 2623.89712487  
Validation score: 0.957034  
Iteration 52, loss = 2595.33693212  
Validation score: 0.957714  
Iteration 53, loss = 2590.92936057  
Validation score: 0.957525  
Iteration 54, loss = 2615.57414957  
Validation score: 0.956998  
Iteration 55, loss = 2628.97969433  
Validation score: 0.957136

Iteration 56, loss = 2574.35117809  
Validation score: 0.955533  
Iteration 57, loss = 2598.76828229  
Validation score: 0.957352  
Iteration 58, loss = 2592.31860233  
Validation score: 0.957899  
Iteration 59, loss = 2629.03775109  
Validation score: 0.957108  
Iteration 60, loss = 2560.21381202  
Validation score: 0.957306  
Iteration 61, loss = 2572.66679465  
Validation score: 0.956254  
Iteration 62, loss = 2580.40615258  
Validation score: 0.957401  
Iteration 63, loss = 2561.93021908  
Validation score: 0.953763  
Iteration 64, loss = 2583.91167570  
Validation score: 0.957855  
Iteration 65, loss = 2577.15733048  
Validation score: 0.954738  
Iteration 66, loss = 2560.55920948  
Validation score: 0.956368  
Iteration 67, loss = 2582.89620406  
Validation score: 0.953563  
Iteration 68, loss = 2539.67821685  
Validation score: 0.954502  
Iteration 69, loss = 2586.50281436  
Validation score: 0.958076  
Iteration 70, loss = 2539.07454030  
Validation score: 0.957650  
Iteration 71, loss = 2539.17665727  
Validation score: 0.957874  
Iteration 72, loss = 2526.09039548  
Validation score: 0.955149  
Iteration 73, loss = 2578.01193607  
Validation score: 0.958080  
Iteration 74, loss = 2568.79862623  
Validation score: 0.957712  
Iteration 75, loss = 2552.06760454  
Validation score: 0.956531  
Iteration 76, loss = 2560.13477279  
Validation score: 0.955206  
Iteration 77, loss = 2482.22410817  
Validation score: 0.957790  
Iteration 78, loss = 2485.20904692  
Validation score: 0.955304  
Iteration 79, loss = 2482.92618286  
Validation score: 0.957447

```

Iteration 80, loss = 2494.84606321
Validation score: 0.955910
Iteration 81, loss = 2508.87559548
Validation score: 0.957711
Iteration 82, loss = 2525.15199528
Validation score: 0.957607
Iteration 83, loss = 2485.24101376
Validation score: 0.958265
Iteration 84, loss = 2475.47789266
Validation score: 0.957764
Iteration 85, loss = 2464.08604742
Validation score: 0.957631
Iteration 86, loss = 2477.27369603
Validation score: 0.957531
Iteration 87, loss = 2472.67163458
Validation score: 0.957814
Iteration 88, loss = 2541.88006435
Validation score: 0.953028
Iteration 89, loss = 2455.09286815
Validation score: 0.956354
Iteration 90, loss = 2451.18673487
Validation score: 0.958016
Iteration 91, loss = 2473.74032244
Validation score: 0.955968
Iteration 92, loss = 2470.03417351
Validation score: 0.957904
Validation score did not improve more than tol=0.000100 for 50 consecutive
epochs. Stopping.

```

	MAE	RMSE	Bias	MAE_pct	RMSE_pct	r2_score
NN						
Train	29.25	71.99	-0.83	17.37	42.77	0.95
Test	35.83	87.70	0.56	17.84	43.66	0.95

```

[240]: X_train, y_train, X_test, y_test = train_test_datasets(df, x_len=12, y_len=1,
    ↪test_loops=12)

```

```

NN = Pipeline([('scaler', MinMaxScaler()),
                ('NN', MLPRegressor(activation='relu',
                                     solver='adam',
                                     early_stopping=True,
                                     n_iter_no_change=50,
                                     validation_fraction=.1,
                                     tol=.0001))])

```

```

param_dist = {'NN__hidden_layer_sizes': [[neuron]*hidden_layer for neuron in
    ↪range(10,60,10) for hidden_layer in range(2,7)],
              'NN__alpha': [5, 1, .5, .1, .05, .01, .001],

```

```

        'NN__learning_rate_init': [.05, .01, .005, .001, .0005],
        'NN__beta_1': [.85, .875, .9, .95, .975, .99, .995],
        'NN__beta_2': [.99, .995, .999, .9995, .9999]}

NN_cv = RandomizedSearchCV(NN, param_dist, cv=10, verbose=2, n_jobs=-1,
    ↪n_iter=200, scoring='neg_mean_absolute_error')
NN_cv.fit(X_train, y_train)

print('Tuned NN Parameters:', NN_cv.best_params_)

```

Fitting 10 folds for each of 200 candidates, totalling 2000 fits  
Tuned NN Parameters: {'learning\_rate\_init': 0.005, 'hidden\_layer\_sizes': [30, 30, 30, 30, 30], 'beta\_2': 0.995, 'beta\_1': 0.975, 'alpha': 0.001}

	MAE	RMSE	Bias	MAE_pct	RMSE_pct	r2_score
NN optimized						
Train	28.55	69.02	-3.51	16.96	41.00	0.96
Test	34.66	86.15	-1.91	17.25	42.89	0.95

```

[469]: X_train, y_train, X_test, y_test = train_test_datasets(df, x_len=12, y_len=1,
    ↪test_loops=20)

NN_final = Pipeline([('scaler', MinMaxScaler()),
    ('NN', MLPRegressor(activation='relu',
        solver='adam',
        early_stopping=True,
        n_iter_no_change=50,
        validation_fraction=.1,
        tol=.0001,
        hidden_layer_sizes=[30, 30, 30, 30,
    ↪30],
        alpha=.001,
        learning_rate_init=.005,
        beta_1=.97,
        beta_2=.995)))]

NN_final.fit(X_train, y_train)

y_train_pred = NN_final.predict(X_train)
y_test_pred = NN_final.predict(X_test)

kpi(y_train, y_train_pred, y_test, y_test_pred, name='NN Final')

```

	MAE	RMSE	Bias	MAE_pct	RMSE_pct	r2_score
NN Final						
Train	29.25	72.58	-1.45	17.93	44.49	0.95
Test	32.10	82.84	-1.93	17.04	43.97	0.95

```
[472]: #Forecasting

X_train, y_train, X_test, y_test = train_test_datasets(df, x_len=12, y_len=1,
↳test_loops=0)

NN_final = Pipeline([('scaler', MinMaxScaler()),
                      ('NN', MLPRegressor(activation='relu',
                                           solver='adam',
                                           early_stopping=True,
                                           n_iter_no_change=50,
                                           validation_fraction=.1,
                                           tol=.0001,
                                           hidden_layer_sizes=[30, 30, 30, 30,
↳30],

                                           alpha=.001,
                                           learning_rate_init=.005,
                                           beta_1=.97,
                                           beta_2=.995)))]

NN_final.fit(X_train, y_train)

forecast = pd.DataFrame(data=NN_final.predict(X_test), index=df.index,
↳columns=['Next_Sales'])
forecast.tail(10)
```

```
[472]:
```

	Next_Sales
Make	
Subaru	237.845909
Suzuki	268.173489
Tata	0.988654
Tazzari	0.988654
Tesla	176.352180
Think	0.988654
Toyota	1151.028659
Volkswagen	1850.869401
Volvo	744.645220
Westfield	0.988654

### 0.3.1 LSTM

```
[511]: X_train, y_train, X_test, y_test = train_test_datasets(df, x_len=3, y_len=1,
↳test_loops=12)

sc = MinMaxScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

n_steps = 3 # =x_len
```

```

n_features = 1

# reshape from [samples, timesteps] into [samples, timesteps, features]
X_train = X_train.reshape((X_train.shape[0], X_train.shape[1], n_features))
X_test = X_test.reshape((X_test.shape[0], X_test.shape[1], n_features))
y_train = y_train.reshape(-1,1)
y_test = y_test.reshape(-1,1)

X_train.shape, X_test.shape, y_train.shape, y_test.shape

```

[511]: ((5720, 3, 1), (1170, 3, 1), (5720, 1), (1170, 1))

```

[512]: tf.keras.backend.clear_session()
tf.random.set_seed(42)
np.random.seed(42)

model = tf.keras.models.Sequential([tf.keras.layers.Conv1D(filters=60,
                                                             kernel_size=5,
                                                             strides=1,
                                                             padding="causal",
                                                             activation="relu",
                                                             ↵
                                                             ↪input_shape=(n_steps, n_features)),
                                     tf.keras.layers.LSTM(60,↵
                                                             ↪return_sequences=True),
                                     tf.keras.layers.LSTM(60), tf.keras.layers.
                                                             ↪Dense(1),
                                     tf.keras.layers.Lambda(lambda x: x * 400)])

model.compile(loss='mean_squared_error', optimizer='adam')

history = model.fit(X_train, y_train, validation_data=(X_test, y_test),↵
                    ↪epochs=100 ,batch_size=16, verbose=0)

```

```

[514]: y_train_pred = model.predict(X_train)
y_test_pred = model.predict(X_test)

kpi(y_train, y_train_pred, y_test, y_test_pred, name='LSTM')

```

	MAE	RMSE	Bias	MAE_pct	RMSE_pct	r2_score
LSTM						
Train	30.29	77.26	-1.80	18.53	47.26	0.94
Test	38.94	106.53	2.26	20.54	56.19	0.92



```
[517]: # Forecasting

X_train, y_train, X_test, y_test = train_test_datasets(df, x_len=3, y_len=1,
↳test_loops=0)

sc = MinMaxScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

n_steps = 3 # =x_len
n_features = 1

# reshape from [samples, timesteps] into [samples, timesteps, features]
X_train = X_train.reshape((X_train.shape[0], X_train.shape[1], n_features))
X_test = X_test.reshape((X_test.shape[0], X_test.shape[1], n_features))
y_train = y_train.reshape(-1,1)
y_test = y_test.reshape(-1,1)
```

```
[501]: tf.keras.backend.clear_session()
tf.random.set_seed(42)
np.random.seed(42)

model1 = tf.keras.models.Sequential([tf.keras.layers.Conv1D(filters=60,
                                                                    kernel_size=5,
                                                                    strides=1,
                                                                    padding="causal",
                                                                    activation="relu",
                                                                    ↳
↳input_shape=(n_steps, n_features)),
                                      tf.keras.layers.LSTM(60,↳
↳return_sequences=True),
                                      tf.keras.layers.LSTM(60), tf.keras.layers.
↳Dense(1),
                                      tf.keras.layers.Lambda(lambda x: x * 400)])

model1.compile(loss='mean_squared_error',optimizer='adam')

history = model1.fit(X_train, y_train, epochs=100 ,batch_size=16, verbose=0)
```

```
[518]: forecast = pd.DataFrame(data=model1.predict(X_test), index=df.index,↳
↳columns=['Next_Sales'])
forecast.head()
```

```
[518]:
```

	Next_Sales
Make	
Alfa Romeo	2.751791
Aston Martin	2.552545
Audi	507.842651
BMW	999.981812
Bentley	2.552545